

Ankita Pandit

Data Engineer | ETL Pipeline Developer | Automation Engineer | 5 Years Experience

Redmond, WA | (716) 603-1325 | ankita.pandit.325@gmail.com | linkedin.com/in/ankitapandit | github.com/ankitapandit

Professional Summary

Experienced Data Engineer with 5 years of industry exposure designing scalable ETL pipeline architectures and automation processes in tech, healthcare and aerospace. Successfully increased Apple's deliverables throughput by 55% and reduced manual efforts by 70% through scripting solutions and Snowflake/Kafka-based pipeline architectures. Currently working on developing AI applications utilizing RAG, CLIP embeddings and LangChain.

Core Competencies

Data Engineering: Snowflake, SQL, Teradata, Airflow, Spark, Scala, Hadoop, Kafka, BigQuery, dbt, Datalake, Databricks, ETL/ELT Pipelines, Data Quality, Data Warehouse, Data Modeling, Data Governance, Metadata Management, DevOps, DataOps, Data Analysis

Technologies: Python, Java, AWS, Bash, CI/CD, Bitbucket, Docker, C++, Tableau, Solr, Splunk, Git, Jenkins, OAuth

AI: OpenAI API, OpenAI CLIP, LangChain, Prompt Engineering, RAG, HuggingFace, GitHub Copilot, Claude Code, OpenAI Codex, Embedding Generation, Vector Search

Web Development: TypeScript, FastAPI, Streamlit, JSON, Next.js, Flask, REST API, HTML, CSS, React.js, JavaScript

Soft Skills: Leadership, Organization, Collaboration, Teamwork, Problem-solving, Initiative, Adaptability, Agile

Work Experience

Career Break | Data Engineering and AI Focus

Jul 2024 - Present | Redmond, WA

- Aimed at skill enhancement via certifications and development of AI applications. (Currently authorized to work for any US/Canada employer, previously US employment permit delayed due to visa change)
- Developed a Food Recommendation Engine using Retrieval-Augmented Generation (RAG) pipeline along with location-based retrieval and user preferences; developed data ingestion and pre-processing workflow for embedding based retrieval. [Python, LangChain, ChromaDB, Streamlit, REST APIs]
- Built LuxeLook AI, which is an AI-based outfit recommendation platform, which makes use of multimodal inputs (uploads by users and CLIP embeddings) to recommend outfits. Developed async **10-item** batch pipelines, wardrobe analytics and event-based styling workflows backed by 135 passing unit tests and 58 benchmark tests; locally benchmarked tag+embed throughput at 10K+ items/s with sub-1ms API p95s in mock environments. [Next.js, FastAPI, Supabase, OpenAI]

Data Engineer | Cognizant Technology Solutions - Apple

Dec 2020 – Jul 2024 | College Station, TX

- Collaborated with cross-functional teams, customers and stakeholders while leading a small team to deliver high-impact solutions for Apple (client) under Cognizant Technology Solutions, increasing deliverables throughput by **55%** within a year.
- Built an end-to-end product showcasing the impact of decommissioning on different stakeholders along with visualizations to depict the impact and BU level dependencies. This helped break down a challenging task of decommissioning into action items and helped in reduction of manual effort by other BUs by **70%** in addition to rationalizing the planning budget and timelines for this new project.
- Re-architected a legacy product performance index application, consolidating processing from **12** to **7** pipelines and shifting job execution to event-driven Kafka-based loading, reducing business-reported escalations from quarterly to near-zero.
- Spearheaded development of multiple ETL pipelines using agile methodology for multiple source ingestion (like Kafka, Oracle DB, Hadoop, MongoDB), maintenance of pipelines and resolve orchestration failures. Ensured high availability and data accuracy; contributing to better performance, reliability for downstream analytics.
- Contributed to GDPR compliance tool within domain applications ensuring alignment with data privacy regulations and improving compliance compatibility score, reducing legal compliance flags by **9%**.
- Developed near real-time data comparison tool with automated discrepancy summary reporting. This increased team's productivity and reducing manual workload by **80%** and **10%** cost reduction using query optimization and consistent reporting.
- Developed DDL migration tool with parser-based pattern recognition and rule-based transformation logic to convert database objects from Teradata to Snowflake. The parsing logic increased the migration efficiency by **3.5x** the manual processes.
- Designed application for validations on RDBMS to cloud migration by streamlining tasks like dependency resolution, request validation, DR snapshot provisioning and pruning, DR health monitoring, migration strategy identification and data quality check using scripting, reducing manual workload by **65%** and increasing throughput and data integrity.

Software Engineer | Cerner Corporation (now Oracle Health)

Mar 2020 – Jun 2020 | Kansas City, MO

- Implemented web application using React.js and Ruby on Rails for a virtual campus tour, enabling new hires to navigate campuses with ease, hence reducing onboarding time by **3** hours per hire.

Software Engineer | Honeywell Technology Solutions Lab

Jul 2017 – Jul 2018 | Bangalore, India

- Developed C++ APIs for flight maintenance and connectivity, incorporating CI/CD pipeline build for latest aircraft applications using Bash that currently supports multiple aircraft models primarily helicopters.
- Configured and deployed RADIUS server for authentication, authorization and accounting (AAA) validation, improving system reliability, security and access control.

Technical Intern | KPMG Global Services

Jan 2017 – May 2017 | Bangalore, India

- Engineered a comprehensive Profit & Loss (P&L) tool and timesheet application using MS VBA which streamlined financial reporting and eliminated outsourcing or licensed software use, driving significant cost savings for the team with no maintenance requirement.

Certifications

- Associate Data Engineer & Data Engineer, DataCamp - 2026
- Data Engineering Foundations in Python, Educative - 2025
- MCP Fundamentals for Building AI Agents, Educative - 2025
- All You Need to Know About Prompt Engineering, Educative - 2025
- Advanced RAG Techniques: Choosing the Right Approach, Educative - 2025
- Master GitHub Copilot, Educative - 2025
- Lean Six Sigma Green Belt Certification, Honeywell Technology Solutions Lab Private Limited - 2018

Education

M.Sc. Computer Science and Engineering | University at Buffalo (SUNY)

Aug 2018 - Feb 2020 | Buffalo, NY

Data Digest – Text extraction and summarization on election violence from news and social media data. [Python, Flask, Frontend, Solr, RESTful API]

Tweet Analyzer - Text search with sentiment analysis, distribution graphs, word cloud, trending hashtags on social media data. [Python, Flask, Frontend, AWS, Solr, RESTful API]